
CS2030 Lecture 9

Java Streams

Henry Chia (hchia@comp.nus.edu.sg)

Semester 2 2022 / 2023

Lecture Outline and Learning Outcomes

- Know how to create **stream** pipelines for *internal* iteration
 - Know the difference between primitive and generic streams
 - Know how to write correct streams that are non-interfering
- Understand **lazy evaluation** in source/intermediate operations, and **eager evaluation** for terminal operations
- Appreciate how lazy evaluation supports **infinite stream**
- Able to implement a basic lazy context by encapsulating a `Supplier` functional interface for *delayed data*
 - Understand how `Lazy::map` can be implemented lazily
- Appreciate the concept of the lambda closure

External Iteration

- An external iteration is defined imperatively
 - e.g. sum of all integers in the closed interval $[1, 10]$

```
jshell> int sum = 0  
sum ==> 0
```

```
jshell> for (int x = 1; x <= 10; x = x + 1) {  
    ...>     sum = sum + x;  
    ...> }
```

```
jshell> sum  
sum ==> 55
```

- Errors could be introduced when
 - sum is initialized wrongly before the loop
 - looping variable x is initialized wrongly
 - loop condition is wrong
 - increment of x is wrong
 - aggregation of sum is wrong

Internal Iteration: Stream

- Internal iteration is defined declaratively
 - e.g. using a primitive integer stream

```
jshell> int sum = IntStream.rangeClosed(1, 10).  
    ...> sum()  
sum ==> 55
```
- Literal meaning “loop through values 1 to 10, and sum them”
- No need to specify how to iterate through elements or use any *mutable* variables — no variable state, no surprises! 😊
- A **stream** is a sequence of elements on which tasks are performed; stream elements move through a sequence of tasks in the stream pipeline
 - E.g. `sum` is assigned with the result of the stream pipeline

Stream Pipeline

- A stream pipeline comprises
 - a **data source**, e.g. `IntStream::rangeClosed` that starts the stream
 - some **intermediate operations**, e.g. `IntStream::map` that specify tasks to perform on a stream's elements

```
jshell> IntStream stream = IntStream.rangeClosed(1, 10).map(x -> x * 2)
stream ==> java.util.stream.IntPipeline$Head@12edcd21
```
 - a **terminal operation**, e.g. `IntStream::sum` that *reduces* the stream elements into a single value

```
jshell> stream.sum()
$.. ==> 110
```
- Each source/intermediate operation returns a new stream of processing steps specified up to that point in the pipeline
- Stream elements within a stream *can only be consumed once*


Reducing a Stream to a Value

- Iterate through `IntStream` elements and reduce to an **int**
int `reduce(int identity, IntBinaryOperator op)`

- `IntBinaryOperator` with single abstract method:
int `applyAsInt(int left, int right)`

```
jshell> IntStream.rangeClosed(1, 10).  
...> reduce(0, (x,y) -> x + y)  
$.. ==> 55
```

```
jshell> IntStream.rangeClosed(1, 10).  
...> reduce(1, (x,y) -> x * y)  
$.. ==> 3628800
```



- Alternative one argument reduce that returns `OptionalInt`
`OptionalInt reduce(IntBinaryOperator op)`

```
jshell> IntStream.range(1, 10).reduce((x, y) -> x < y ? x : y)  
$.. ==> OptionalInt[9]
```

```
jshell> IntStream.range(1, 1).reduce((x, y) -> x < y ? x : y)  
$.. ==> OptionalInt.empty
```

flatMap Method in Stream

- How about nested loops?

```
for (x = 1; x <= 3; x++)  
    for (y = x; y <= 3; y++)  
        System.out.println((x * y) + " "); // output is 1 2 3 4 6 9
```

- map tries to map each stream element into one other stream

```
jshell> IntStream.of(1,2,3).  
...> map(x -> IntStream.rangeClosed(x,3).map(y -> x * y))  
| Error:  
| incompatible types: bad return type in lambda expression  
|     java.util.stream.IntStream cannot be converted to int  
| map(x -> IntStream.rangeClosed(x,3).map(y -> x * y))  
|     ^-----^
```

- flatMap transforms each stream element into a stream of other elements (either zero or more) by taking in a function that produces another stream, and then *flattens* it

```
jshell> IntStream.of(1,2,3).  
...> flatMap(x -> IntStream.rangeClosed(x,3).map(y -> x * y)).  
...> forEach(x -> System.out.print(x + " "))  
1 2 3 4 6 9
```

Generic Stream<T>

- Stream<T> is a stream over reference-typed objects, e.g.

```
jshell> int sum = Stream.<Integer>iterate(1, x -> x <= 10, x -> x + 1).  
    ...> reduce(0, (x,y) -> x + y)  
sum ==> 55
```

- boxed() wraps stream elements in its wrapper type

```
jshell> Stream.<Integer> stream = IntStream.rangeClosed(1, 10).boxed()  
stream ==> java.util.stream.IntPipeline$1@5010be6  
  
jshell> List.<Integer> list = stream.toList()  
list ==> [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
```

- mapToObj converts from primitive to generic stream

```
jshell> IntStream.rangeClosed(1, 10).  
    ...> mapToObj(x -> "<" + x + ">").  
    ...> toList()  
$.. ==> [<1>, <2>, <3>, <4>, <5>, <6>, <7>, <8>, <9>, <10>]
```

- Stream::toList() converts generic stream to generic list
- List::stream() converts generic list to generic stream

Correctness of Streams

- To ensure correct execution, stream operations

- must not interfere with stream data

```
jshell> List<String> list = new ArrayList<String>(
...>     List.of("abc", "def", "xyz"))
list ==> [abc, def, xyz]

jshell> list.stream().peek(str -> {
...>     if (str.equals("xyz")) { list.add("pqr"); }
...> }).forEach(x -> {})
| Exception java.util.ConcurrentModificationException
| ...
```

- *preferably* stateless with no side effects

- ▷ operations like `filter` and `map` are **stateless**, i.e. processing one stream element does not depend on other stream elements
- ▷ **stateful** operations like `sorted`, `limit`, `distinct`, etc. depend on the current state

Lazy Evaluation in Streams

- Source/intermediate operations use **lazy evaluation**
 - does not perform any operations on stream's elements until a terminal operation is called
- Terminal operations use **eager evaluation**
 - performs the requested operation as soon as it is called

```
jshell> Stream.<Integer>iterate(1, x -> x + 1).  
...> limit(5).  
...> peek(x -> System.out.println("limit: " + x)).  
...> filter(x -> x % 2 == 0).  
...> peek(x -> System.out.println("filter: " + x)).  
...> map(x -> 2 * x).  
...> peek(x -> System.out.println("map: " + x)).  
...> sum()
```

```
limit: 1  
limit: 2  
filter: 2  
map: 4  
limit: 3  
limit: 4  
filter: 4  
map: 8  
limit: 5  
$.. ==> 12
```

Infinite Stream

- Lazy evaluation allows us to work with infinite streams that represent an infinite number of elements
 - `Stream<T>::generate(Supplier<T> supplier)` produces an infinite sequence of values generated by `supplier`
 - `Stream<T>::iterate(T seed, UnaryOperator<T> next)` produces an infinite sequence by repeatedly applying the function `next` starting with the `seed` value
- Intermediate operations, e.g. `limit`, can be used to restrict the total number of elements in the stream

```
jshell> Stream.<Integer>iterate(1, x -> x + 1).  
...> filter(x -> x % 2 == 1).  
...> limit(20). // find first 20 odd numbers  
...> forEach(x -> System.out.print(x + " "))  
1 3 5 7 9 11 13 15 17 19 21 23 25 27 29 31 33 35 37 39
```

Lazy Class

- To understand how lazy evaluation works, define a Lazy class

```
import java.util.function.Supplier;

class Lazy<T> implements Supplier<T> {
    private final Supplier<T> supplier;

    private Lazy(Supplier<T> supplier) {
        this.supplier = supplier;
    }

    static <T> Lazy<T> of(Supplier<T> supplier) {
        return new Lazy<T>(supplier);
    }

    static <T> Lazy<T> of(T t) {
        return new Lazy<T>(() -> t);
    }

    @Override
    public T get() {
        return supplier.get();
    }
}
```

```
jshell> int foo() {
...>     System.out.println("foo");
...>     return -1;
...> }
| created method foo()
jshell> Lazy<Integer> lazy = Lazy.of(foo())
foo
$.. ==> Lazy@ae45eb6
jshell> lazy.get()
$.. ==> -1
jshell> lazy = Lazy.<Integer>of(() -> foo())
$.. ==> Lazy@6f7fd0e6
jshell> lazy.get()
foo
$.. ==> -1
```

- `Lazy.of(foo())` evaluates `foo` method *eagerly*
- `Lazy.of(() -> foo())` evaluates `foo` *lazily*, i.e. only when `get()` is invoked sometime later

Mapping a Lazy Value

- Define map that returns a new Lazy

```
<R> Lazy<R> map(Function<? super T, ? extends R> mapper) {  
    Supplier<R> supplier = () -> mapper.apply(this.get());  
    return Lazy.<R>of(supplier);  
}
```

```
jshell> Lazy<Integer> i = Lazy.<String>of(() -> "abc").  
...>    map(x -> { System.out.println("map1"); return x.length(); }).  
...>    map(x -> { System.out.println("map2"); return x * 2; })  
i ==> Lazy@51565ec2 // map is not evaluated until a get()  
  
jshell> i.get() // map is lazily evaluated :)  
map1  
map2  
$.. ==> 6
```

- What about the following implementation of map?

```
<R> Lazy<R> map(Function<? super T, ? extends R> mapper) {  
    R r = mapper.apply(this.get());  
    return Lazy.<R>of(() -> r);  
}
```

Lambda Closure

- Lambdas declared inside a method are also *local classes*

```
jshell> class A {  
...>     private final int x;  
...>     A(int x) {  
...>         this.x = x;  
...>     }  
...>     Function<Integer,Integer> f(int y) {  
...>         return z -> this.x + y + z; // or A.this.x + y + z ?  
...>     }  
...> }
```

```
| created class A
```

```
jshell> Function<Integer,Integer> fn = new A(1).f(2)  
fn ==> A$$Lambda$14/1196765369@26be92ad
```

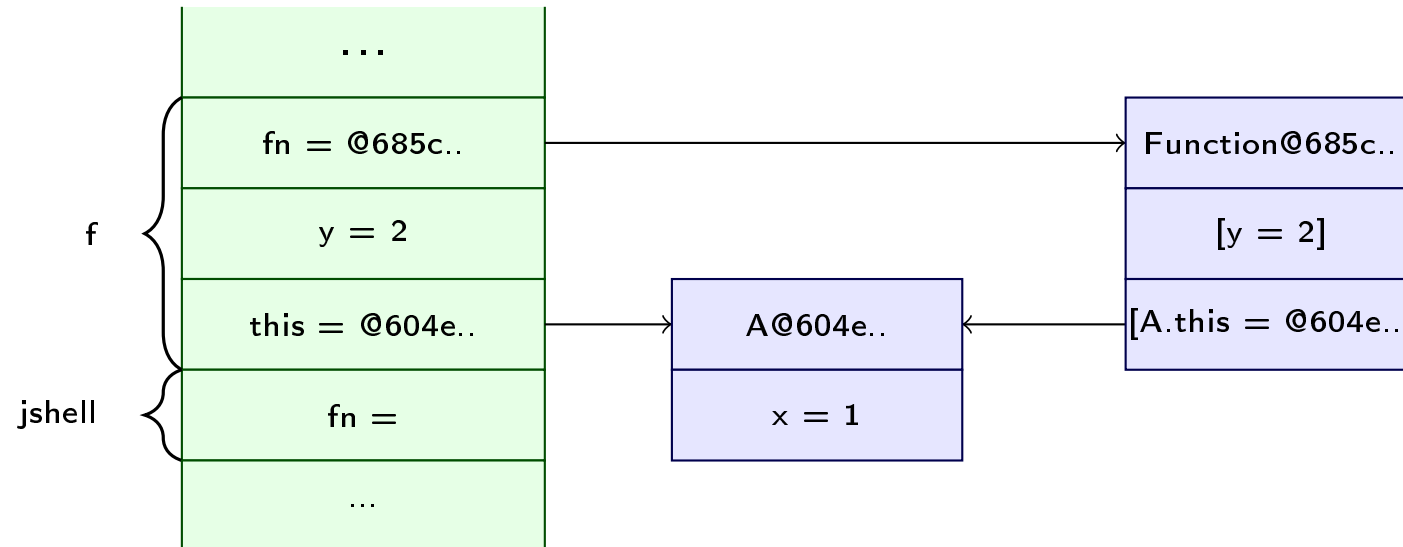
```
jshell> fn.apply(3)  
$.. ==> 6
```

- *Lambda closure*: lambda expression closes over it's enclosing method and class
 - captures the variables of the enclosing method and reference to the enclosing class

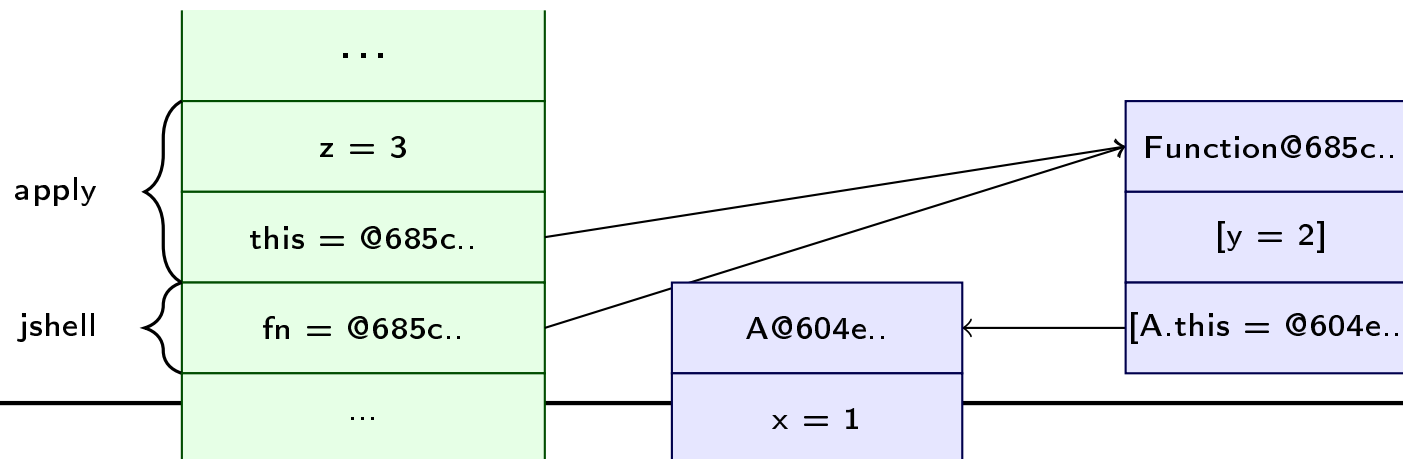
Java Memory Model

- Memory model just before returning from the method `f`

```
jshell> Function<Integer,Integer> fn = new A(1).f(2)
```



- Memory model upon invoking the method `fn.apply(3)`



Exercise

- Let's repeat the exercise in the previous lecture but with method `f` returning a lambda expression instead

```
jshell> class A {  
...>     Integer apply(int x) {  
...>         return x * 10;  
...>     }  
...>     Function<Integer,Integer> f(int y) {  
...>         return z -> A.this.apply(z) + y;  
...>     }  
...> }  
| modified class A
```

- What is the outcome of `new A().f(2).apply(3)`?
- Now replace `A.this.apply(z)` in method `foo` with `this.apply(z)`. Does it compile?
 - what is the outcome of `new A().f(2).apply(3)` now?
 - what is the difference as compared to returning an anonymous inner class?